

Savanna

Generated by Doxygen 1.8.13

Contents

1	Namespace Index	1
1.1	Namespace List	1
2	Hierarchical Index	3
2.1	Class Hierarchy	3
3	Class Index	5
3.1	Class List	5
4	Namespace Documentation	7
4.1	codar.savanna Namespace Reference	7
4.1.1	Detailed Description	7
4.2	codar.savanna.consumer Namespace Reference	7
4.2.1	Detailed Description	8
4.3	codar.savanna.exc Namespace Reference	8
4.3.1	Detailed Description	8
4.4	codar.savanna.machines Namespace Reference	8
4.4.1	Detailed Description	8
4.4.2	Variable Documentation	8
4.4.2.1	cori	9
4.4.2.2	summit	9
4.4.2.3	theta	9
4.4.2.4	titan	9
4.5	codar.savanna.main Namespace Reference	10
4.5.1	Detailed Description	10
4.6	codar.savanna.model Namespace Reference	10
4.6.1	Detailed Description	10
4.7	codar.savanna.producer Namespace Reference	11
4.7.1	Detailed Description	11
4.8	codar.savanna.scheduler Namespace Reference	11
4.8.1	Detailed Description	11
4.9	codar.savanna.status Namespace Reference	11
4.9.1	Detailed Description	11

5	Class Documentation	13
5.1	codar.savanna.scheduler.JobList Class Reference	13
5.1.1	Detailed Description	14
5.1.2	Member Function Documentation	14
5.1.2.1	pop_job()	14
5.2	codar.savanna.producer.JSONFilePipelineReader Class Reference	15
5.2.1	Detailed Description	15
5.3	codar.savanna.machines.Machine Class Reference	16
5.3.1	Detailed Description	17
5.3.2	Member Function Documentation	17
5.3.2.1	get_scheduler_options()	17
5.4	codar.savanna.machines.MachineNode Class Reference	17
5.5	codar.savanna.exc.MachineNotFound Class Reference	18
5.6	codar.savanna.runners.MPIRunner Class Reference	19
5.7	codar.savanna.model.NodeConfig Class Reference	20
5.7.1	Constructor & Destructor Documentation	20
5.7.1.1	__init__()	20
5.8	codar.savanna.node_layout.NodeLayout Class Reference	21
5.8.1	Detailed Description	22
5.8.2	Member Function Documentation	22
5.8.2.1	add_node()	22
5.8.2.2	default_no_share_layout()	22
5.8.2.3	get_node_containing_code()	23
5.8.2.4	group_codes_by_node()	23
5.8.2.5	serialize_to_dict()	23
5.8.2.6	validate()	23
5.9	codar.savanna.model.Pipeline Class Reference	24
5.9.1	Member Function Documentation	25
5.9.1.1	force_kill_all()	25
5.9.1.2	from_data()	25

5.9.1.3	set_ppn()	26
5.9.1.4	set_total_nodes()	26
5.10	codar.savanna.consumer.PipelineRunner Class Reference	26
5.10.1	Detailed Description	27
5.10.2	Member Function Documentation	28
5.10.2.1	kill_all()	28
5.10.2.2	pipeline_finished()	28
5.10.2.3	run_finished()	28
5.10.2.4	run_pipelines()	28
5.10.2.5	stop()	29
5.11	codar.savanna.status.PipelineState Class Reference	29
5.12	codar.savanna.model.Run Class Reference	30
5.12.1	Detailed Description	32
5.12.2	Member Function Documentation	32
5.12.2.1	add_callback()	32
5.12.2.2	exception()	32
5.12.2.3	from_data()	32
5.12.2.4	get_nodes_used()	33
5.12.2.5	kill()	33
5.12.2.6	killed()	33
5.12.2.7	succeeded()	33
5.12.2.8	timed_out()	34
5.13	codar.savanna.runners.Runner Class Reference	34
5.14	codar.savanna.exc.SavannaException Class Reference	35
5.15	codar.savanna.machines.SummitNode Class Reference	36
5.15.1	Member Function Documentation	36
5.15.1.1	validate_layout()	37
5.16	codar.savanna.runners.SummitRunner Class Reference	37
5.16.1	Member Function Documentation	38
5.16.1.1	wrap_deprecated()	38
5.17	codar.savanna.status.WorkflowStatus Class Reference	39

Chapter 1

Namespace Index

1.1 Namespace List

Here is a list of all documented namespaces with brief descriptions:

codar.savanna	7
codar.savanna.consumer	7
codar.savanna.exc	8
codar.savanna.machines	8
codar.savanna.main	10
codar.savanna.model	10
codar.savanna.producer	11
codar.savanna.scheduler	11
codar.savanna.status	11

Chapter 2

Hierarchical Index

2.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

Exception	
codar.savanna.exc.SavannaException	35
codar.savanna.exc.MachineNotFound	18
codar.savanna.machines.MachineNode	17
codar.savanna.machines.SummitNode	36
codar.savanna.model.NodeConfig	20
object	
codar.savanna.consumer.PipelineRunner	26
codar.savanna.machines.Machine	16
codar.savanna.model.Pipeline	24
codar.savanna.node_layout.NodeLayout	21
codar.savanna.producer.JSONFilePipelineReader	15
codar.savanna.runners.Runner	34
codar.savanna.runners.MPIRunner	19
codar.savanna.runners.SummitRunner	37
codar.savanna.scheduler.JobList	13
codar.savanna.status.PipelineState	29
Thread	
codar.savanna.model.Run	30
codar.savanna.status.WorkflowStatus	39

Chapter 3

Class Index

3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

codar.savanna.scheduler.JobList	13
codar.savanna.producer.JSONFilePipelineReader	15
codar.savanna.machines.Machine	16
codar.savanna.machines.MachineNode	17
codar.savanna.exc.MachineNotFound	18
codar.savanna.runners.MPIRunner	19
codar.savanna.model.NodeConfig	20
codar.savanna.node_layout.NodeLayout	21
codar.savanna.model.Pipeline	24
codar.savanna.consumer.PipelineRunner	26
codar.savanna.status.PipelineState	29
codar.savanna.model.Run	30
codar.savanna.runners.Runner	34
codar.savanna.exc.SavannaException	35
codar.savanna.machines.SummitNode	36
codar.savanna.runners.SummitRunner	37
codar.savanna.status.WorkflowStatus	39

Chapter 4

Namespace Documentation

4.1 `codar.savanna` Namespace Reference

Namespaces

- [consumer](#)
- [exc](#)
- [machines](#)
- [main](#)
- [model](#)
- [producer](#)
- [scheduler](#)
- [status](#)

4.1.1 Detailed Description

Classes for running pipelines of MPI tasks based on a specified total process limit. The system is designed to use two + N threads:

1. consumer thread: get pipelines from queue and execute them when process slots become available. Stops when a None pipeline is received.
2. producer thread: add pipelines to queue. Can be from file or from network service.
3. monitor threads: each process spawned by the consumer thread has a monitor thread that blocks on the processes completing with a timeout, and kills the process if it's not done after the timeout is reached.

4.2 `codar.savanna.consumer` Namespace Reference

Classes

- class [PipelineRunner](#)

4.2.1 Detailed Description

Classes for 'consuming' pipelines - running groups of MPI tasks based on a specified total process limit.

4.3 `codar.savanna.exc` Namespace Reference

Classes

- class [MachineNotFound](#)
- class [SavannaException](#)

4.3.1 Detailed Description

Exceptions.

4.4 `codar.savanna.machines` Namespace Reference

Classes

- class [Machine](#)
- class [MachineNode](#)
- class [SummitNode](#)

Functions

- def `get_by_name` (name)

Variables

- **SCHEDULER_OPTIONS** = set(["project", "queue", "constraint", "license"])
- **local** = [Machine](#)("local", "local", "mpiexec", MachineNode, processes_per_node=1)
- **titan**
- **cori**
- **theta**
- **summit**

4.4.1 Detailed Description

Configuration for machines supported by Codar.

4.4.2 Variable Documentation

4.4.2.1 cori

codar.savanna.machines.cori

Initial value:

```
1 = Machine('cori', "slurm", "srun", MachineNode,
2           processes_per_node=32, node_exclusive=True,
3           dataspace_servers_per_node=4,
4           scheduler_options=dict(project="",
5                                   queue="debug",
6                                   constraint="haswell",
7                                   license="SCRATCH,project"))
```

4.4.2.2 summit

codar.savanna.machines.summit

Initial value:

```
1 = Machine('summit', "ibm_lsf", "jsrun", SummitNode,
2           processes_per_node=42, node_exclusive=True,
3           scheduler_options=dict(project=""))
```

4.4.2.3 theta

codar.savanna.machines.theta

Initial value:

```
1 = Machine('theta', "cobalt", "aprun", MachineNode,
2           processes_per_node=64, node_exclusive=True,
3           dataspace_servers_per_node=8,
4           scheduler_options=dict(project="",
5                                   queue="debug-flat-quad"))
```

4.4.2.4 titan

codar.savanna.machines.titan

Initial value:

```
1 = Machine('titan', "pbs", "aprun", MachineNode,
2           processes_per_node=16, node_exclusive=True,
3           scheduler_options=dict(project="", queue="debug"),
4           dataspace_servers_per_node=4)
```

4.5 codar.savanna.main Namespace Reference

Functions

- def `parse_args ()`
- def `main ()`
- def `get_job_id ()`

Variables

- `consumer = None`

4.5.1 Detailed Description

Main program for executing workflow script with different producers and runners.

4.6 codar.savanna.model Namespace Reference

Classes

- class [NodeConfig](#)
- class [Pipeline](#)
- class [Run](#)

Variables

- string `STDOUT_NAME` = 'codar.workflow.stdout'
- string `STDERR_NAME` = 'codar.workflow.stderr'
- string `RETURN_NAME` = 'codar.workflow.return'
- string `WALLTIME_NAME` = 'codar.workflow.walltime'
- int `KILL_WAIT` = 30
- int `WAIT_DELAY_KILL` = 30
- int `WAIT_DELAY_GIVE_UP` = 120

4.6.1 Detailed Description

Classes for tracking pipelines and the runs within each pipeline in separate monitor threads that synchronize state.

Note that there is state tracked in these classes which is not available just by looking at the return code. In particular, a run may be killed for several different reasons: external signal, run timeout reached, other run in pipeline failed (when kill on partial fail is set), or if the entire workflow is killed.

The goal here is to provide as much information as possible about why a pipeline failed, to make an informed decision about whether it is worth running again when the workflow is restarted, or if it's failure was more permanent and not subject to outside forces like the job walltime expiring.

4.7 codar.savanna.producer Namespace Reference

Classes

- class [JSONFilePipelineReader](#)

4.7.1 Detailed Description

Classes for producing pipelines.

4.8 codar.savanna.scheduler Namespace Reference

Classes

- class [JobList](#)

4.8.1 Detailed Description

Classes related to finding a job that can run on available resources. Does not assume any knowledge of how long each job will take. Designed for greedy search of a job that will fit whenever resources are freed.

In the context of Cheetah workflows, it's unlikely that there will be more than a few hundred jobs, so it's not worth optimizing the python search code very much. It is however worth making sure that a job is run when resources are available, since super computer resources are expensive. Basically it's worth doing some work in python to make sure we start a big unit of work on compute nodes.

4.9 codar.savanna.status Namespace Reference

Classes

- class [PipelineState](#)
- class [WorkflowStatus](#)

Variables

- string **NOT_STARTED** = 'not_started'
- string **RUNNING** = 'running'
- string **DONE** = 'done'
- string **KILLED** = 'killed'
- string **REASON_TIMEOUT** = 'timeout'
- string **REASON_FAILED** = 'failed'
- string **REASON_SUCCEEDED** = 'succeeded'
- string **REASON_EXCEPTION** = 'exception'
- string **REASON_NOFIT** = 'nofit'

4.9.1 Detailed Description


Class for maintaining state of all FOB runs that the workflow consumer is managing. State is saved in a JSON file, overwritten on each state change.

Chapter 5


Class Documentation

5.1 `codar.savanna.scheduler.JobList` Class Reference

Inheritance diagram for `codar.savanna.scheduler.JobList`:



Collaboration diagram for `codar.savanna.scheduler.JobList`:



Public Member Functions

- `def __init__ (self, costfn, initial_jobs=None)`
- `def add_job (self, job)`
- `def pop_job (self, max_cost)`
- `def __len__ (self)`

5.1.1 Detailed Description

Manage a job list that can find and remove the highest cost job that doesn't exceed `max_cost` and insert new jobs.

The job objects can be any type, but a key function must be provided that takes an instance of a job and returns it's cost.

Uses a coordinated pair of sort list for costs and jobs, along with the `bisect` module. A linked list might be more efficient, since the list copy on insert and delete may dominate the time to do a linear search of a small list, but it's likely fine either way for the sizes we will encounter.

5.1.2 Member Function Documentation

5.1.2.1 `pop_job()`

```
def codar.savanna.scheduler.JobList.pop_job (
    self,
    max_cost )
```


Get the highest cost job that doesn't exceed `max_cost`, and remove it from the job list. Raises `IndexError` if the job list is empty, returns `None` if no suitable jobs exist in the list.

The documentation for this class was generated from the following file:


- `scheduler.py`

5.2 codar.savanna.producer.JSONFilePipelineReader Class Reference

Inheritance diagram for codar.savanna.producer.JSONFilePipelineReader:



Collaboration diagram for codar.savanna.producer.JSONFilePipelineReader:



Public Member Functions

- `def __init__(self, file_path)`
- `def read_pipelines(self)`

Public Attributes

- `file_path`

5.2.1 Detailed Description


Load pipelines from a file formatted as a new line separated list of JSON documents. Each JSON document must be a list containing dictionaries, each dictionary describing a code to run as part of the pipeline.

The documentation for this class was generated from the following file:


- `producer.py`

5.3 `codar.savanna.machines.Machine` Class Reference

Inheritance diagram for `codar.savanna.machines.Machine`:



Collaboration diagram for `codar.savanna.machines.Machine`:



Public Member Functions

- `def __init__(self, name, scheduler_name, runner_name, node_class, processes_per_node=None, node_exclusive=False, scheduler_options=None, dataspace_servers_per_node=1)`
- `def get_scheduler_options(self, options)`
- `def get_nodes_reqd(self)`

Public Attributes

- `name`
- `scheduler_name`
- `runner_name`
- `node_class`
- `processes_per_node`
- `node_exclusive`
- `scheduler_options`
- `dataspace_servers_per_node`

5.3.1 Detailed Description

Class to represent configuration of a specific Supercomputer or workstation, including the scheduler and runner used by the machine. This can be used to map an experiment to run on the machine without having to define machine specific parameter for every experiment separately.

5.3.2 Member Function Documentation

5.3.2.1 get_scheduler_options()

```
def codar.savanna.machines.Machine.get_scheduler_options (
    self,
    options )
```


Validate supplied options and add default values where missing.
Returns a new dictionary.

The documentation for this class was generated from the following file:

- machines.py

5.4 codar.savanna.machines.MachineNode Class Reference

Inheritance diagram for codar.savanna.machines.MachineNode:



Public Member Functions

- def **__init__** (self, num_cpus, num_gpus)
- def **validate_layout** (self)
- def **to_json** (self)

Public Attributes


- **cpu**
- **gpu**

The documentation for this class was generated from the following file:


- machines.py

5.5 `codar.savanna.exc.MachineNotFound` Class Reference

Inheritance diagram for `codar.savanna.exc.MachineNotFound`:



Collaboration diagram for `codar.savanna.exc.MachineNotFound`:



Public Member Functions


- `def __init__(self, machine_name)`

The documentation for this class was generated from the following file:


- `exc.py`

5.6 `codar.savanna.runners.MPIRunner` Class Reference

Inheritance diagram for `codar.savanna.runners.MPIRunner`:



Collaboration diagram for `codar.savanna.runners.MPIRunner`:



Public Member Functions

- def **__init__** (self, exe, nprocs_arg, nodes_arg=None, tasks_per_node_arg=None, hostfile=None)
- def **wrap** (self, run, sched_args, find_in_path=True)

Public Attributes

- **exe**
- **nprocs_arg**
- **nodes_arg**
- **tasks_per_node_arg**
- **hostfile**

The documentation for this class was generated from the following file:

- runners.py

5.7 codar.savanna.model.NodeConfig Class Reference

Public Member Functions

- def **__init__** (self)

Public Attributes

- **num_ranks_per_node**
- **cpu**
- **gpu**

5.7.1 Constructor & Destructor Documentation

5.7.1.1 __init__()

```
def codar.savanna.model.NodeConfig.__init__ (
    self )
```


Intended to look like
 cpu = [0=[], 1=[], 2=[], 3=[]]
 gpu = [0=[], 1=[], 2=[], 3=[]]

The documentation for this class was generated from the following file:


- model.py

5.8 `codar.savanna.node_layout.NodeLayout` Class Reference

Inheritance diagram for `codar.savanna.node_layout.NodeLayout`:



Collaboration diagram for `codar.savanna.node_layout.NodeLayout`:



Public Member Functions

- `def __init__ (self, layout_list)`
- `def add_node (self, node_dict)`
- `def get_node_containing_code (self, code)`
- `def codes_per_node (self)`
- `def shared_nodes (self)`
- `def ppn (self)`
- `def validate (self, ppn, codes_per_node, shared_nodes)`
- `def as_data_list (self)`
- `def serialize_to_dict (self)`
- `def copy (self)`
- `def group_codes_by_node (self)`
- `def populate_remaining (self, rc_names, ppn)`
- `def default_no_share_layout (cls, ppn, code_names)`

Public Attributes

- `layout_list`
- `layout_map`

5.8.1 Detailed Description

Class representing options on how to organize a multi-exe task across many nodes. It is the scheduler model's job to take this and produce the correct scheduler and runner options to make this happen, or raise an error if it's not possible. Note that this will generally be different for each machine unless it is very simple and supported uniformly by all desired machines.

A layout is represented as a list of dictionaries, where each dictionary described codes to be run together on a single node. The keys are the names of the codes, and the values are the number of processes to assign to each.

5.8.2 Member Function Documentation

5.8.2.1 `add_node()`

```
def codar.savanna.node_layout.NodeLayout.add_node (
    self,
    node_dict )
```

Add a node to an existing layout, e.g. add sosflow.

5.8.2.2 `default_no_share_layout()`

```
def codar.savanna.node_layout.NodeLayout.default_no_share_layout (
    cls,
    ppn,
    code_names )
```

Create a layout object for the specified codes and ppn, where each code uses max procs on it's own node.

5.8.2.3 get_node_containing_code()

```
def codar.savanna.node_layout.NodeLayout.get_node_containing_code (
    self,
    code )
```

Get node dict containing the specified code. Raises KeyError if not found.

5.8.2.4 group_codes_by_node()

```
def codar.savanna.node_layout.NodeLayout.group_codes_by_node (
    self )
```

Return a list of dicts, where each list represents codes on a node, and a dict key for ppn
Example: [{sim,analysis1}, {analysis2}, {viz}].
Must take Summit NodeConfigs into account

5.8.2.5 serialize_to_dict()

```
def codar.savanna.node_layout.NodeLayout.serialize_to_dict (
    self )
```

Get a copy of the data list passed to the constructor, suitable for JSON serialization.

5.8.2.6 validate()

```
def codar.savanna.node_layout.NodeLayout.validate (
    self,
    ppn,
    codes_per_node,
    shared_nodes )
```


Given a machine ppn and max number of codes (e.g. 4 on cori), raise a ValueError if the specified layout won't fit.
Dont modify this yet, this is being used by the tests

The documentation for this class was generated from the following file:


- node_layout.py

5.9 codar.savanna.model.Pipeline Class Reference

Inheritance diagram for codar.savanna.model.Pipeline:



Collaboration diagram for codar.savanna.model.Pipeline:



Public Member Functions

- def **__init__** (self, pipe_id, runs, working_dir, total_nodes, machine_name, kill_on_partial_failure=False, post_process_script=None, post_process_args=None, post_process_stop_on_failure=False, node_layout=None, launch_mode=None)
- def **from_data** (cls, data)
- def **start** (self, consumer, nodes_assigned, runner=None)
- def **run_finished** (self, run)
- def **run_post_process_script** (self)
- def **add_done_callback** (self, fn)
- def **remove_done_callback** (self, fn)
- def **add_fatal_callback** (self, fn)
- def **remove_fatal_callback** (self, fn)
- def **get_nodes_used** (self)
- def **set_ppn** (self, ppn)
- def **set_total_nodes** (self)
- def **get_state** (self)
- def **get_pids** (self)
- def **force_kill_all** (self)
- def **join_all** (self)

Public Attributes

- **id**
- **runs**
- **working_dir**
- **kill_on_partial_failure**
- **post_process_script**
- **post_process_args**
- **post_process_stop_on_failure**
- **node_layout**
- **machine_name**
- **done_callbacks**
- **fatal_callbacks**
- **total_procs**
- **log_prefix**
- **total_nodes**
- **launch_mode**
- **nodes_assigned**

5.9.1 Member Function Documentation

5.9.1.1 `force_kill_all()`

```
def codar.savanna.model.Pipeline.force_kill_all (
    self )
```

Kill all runs and don't run post processing. Note that this call may block waiting for all runs to be started, to avoid confusing races. If the pipeline is already done, this does nothing. If one or more runs are still active, or have not yet been marked as finished, then it will mark the entire pipeline as killed so it can be re-run from scratch on a restart if desired.

5.9.1.2 `from_data()`

```
def codar.savanna.model.Pipeline.from_data (
    cls,
    data )
```

Create Pipeline instance from dictionary data structure, containing at least "id" and "runs" keys. The "runs" key must have a list of dict, and each dict is parsed using `Run.from_data`. Raises `KeyError` if a required key is missing.

5.9.1.3 set_ppn()

```
def codar.savanna.model.Pipeline.set_ppn (
    self,
    ppn )
```

Determine number of nodes needed to run pipeline with the specified node layout or full occupancy layout with ppn. Also updates runs to set node and task per node counts.
TODO: This should be set by Cheetah in fobs.json

5.9.1.4 set_total_nodes()

```
def codar.savanna.model.Pipeline.set_total_nodes (
    self )
```


To be deprecated

The documentation for this class was generated from the following file:


- model.py

5.10 codar.savanna.consumer.PipelineRunner Class Reference

Inheritance diagram for codar.savanna.consumer.PipelineRunner:



Collaboration diagram for codar.savanna.consumer.PipelineRunner:



Public Member Functions

- def **__init__** (self, runner, max_nodes, machine_name, processes_per_node, status_file=None)
- def **add_pipeline** (self, p)
- def **stop** (self)
- def **kill_all** (self)
- def **run_finished** (self, run)
- def **pipeline_finished** (self, pipeline)
- def **pipeline_fatal** (self, pipeline)
- def **run_pipelines** (self)

Public Attributes

- **max_nodes**
- **machine_name**
- **ppn**
- **runner**
- **job_list_cv**
- **job_list**
- **free_cv**
- **free_nodes**
- **pipelines_lock**
- **pipelines**
- **allocated_nodes**

5.10.1 Detailed Description

Runner that assumes a homogenous set of nodes. Now only support only node based limiting (although process limiting can be emulated by setting process_per_node=1 and max_nodes=max_procs).

Threading model: assumes there could be multiple producer threads calling add_pipeline, e.g. if using a dynamic job submission model based on results of previous jobs. Pipelines and each Run in a pipeline are all executed in separate threads, so their notification callbacks execute in separate threads, and their threads must be joined before exiting. The stop and kill_all methods could be called from any of the producer, Pipeline or Run threads.

5.10.2 Member Function Documentation

5.10.2.1 kill_all()

```
def codar.savanna.consumer.PipelineRunner.kill_all (
    self )
```

Kill all running processes spawned by this consumer and don't start any new processes.

5.10.2.2 pipeline_finished()

```
def codar.savanna.consumer.PipelineRunner.pipeline_finished (
    self,
    pipeline )
```

Monitor thread(s) should call this as pipelines complete.

5.10.2.3 run_finished()

```
def codar.savanna.consumer.PipelineRunner.run_finished (
    self,
    run )
```

TO BE DEPRECATED.

Monitor thread(s) should call this as runs complete. To be deprecated, as the functionality fails when `node_layout` is set to `node-sharing`.

This means that for `node_exclusive`, resources held by a run are not released when the run terminates. For `kill_on_partial_failure=False`, this could lead to unused resources, which is ok.

5.10.2.4 run_pipelines()

```
def codar.savanna.consumer.PipelineRunner.run_pipelines (
    self )
```

Main loop of consumer thread. Does not return until all child threads are complete.

5.10.2.5 stop()

```
def codar.savanna.consumer.PipelineRunner.stop (
    self )
```


Signal to stop when all pipelines are finished. Don't allow adding new pipelines.

The documentation for this class was generated from the following file:


- consumer.py

5.11 codar.savanna.status.PipelineState Class Reference

Inheritance diagram for codar.savanna.status.PipelineState:



Collaboration diagram for codar.savanna.status.PipelineState:



Public Member Functions

- `def __init__ (self, pipeline_id, state, reason=None, return_codes=None)`
- `def as_data (self)`

Public Attributes


- `id`
- `state`
- `reason`
- `return_codes`

The documentation for this class was generated from the following file:


- `status.py`

5.12 `codar.savanna.model.Run` Class Reference

Inheritance diagram for `codar.savanna.model.Run`:



Collaboration diagram for `codar.savanna.model.Run`:



Public Member Functions

- def **__init__** (self, name, exe, args, sched_args, env, working_dir, timeout=None, nprocs=1, res_set=None, stdout_path=None, stderr_path=None, return_path=None, walltime_path=None, log_prefix=None, sleep_after=None, depends_on_runs=None, hostfile=None, runner_override=False)
- def **from_data** (cls, data)
- def **mpmd_run** (cls, runs)
- def **set_runner** (self, runner)
- def **timed_out** (self)
- def **killed** (self)
- def **exception** (self)
- def **succeeded** (self)
- def **add_callback** (self, fn)
- def **remove_callback** (self, fn)
- def **run** (self)
- def **kill** (self)
- def **get_returncode** (self)
- def **get_pid** (self)
- def **close** (self)
- def **join** (self)
- def **get_nodes_used** (self)
- def **create_node_config** (self)

Public Attributes

- **name**
- **exe**
- **args**
- **sched_args**
- **env**
- **working_dir**
- **timeout**
- **nprocs**
- **res_set**
- **stdout_path**
- **stderr_path**
- **return_path**
- **walltime_path**
- **sleep_after**
- **log_prefix**
- **runner**
- **callbacks**
- **nodes**
- **tasks_per_node**
- **depends_on_runs**
- **hostfile**
- **machine**
- **nodes_assigned**
- **node_config**
- **erf_file**
- **runner_override**

5.12.1 Detailed Description

Manage running a single executable within a pipeline. When start is called, it will launch the process with Popen and call wait in the new thread with a timeout, killing if the process does not finish in time.

5.12.2 Member Function Documentation

5.12.2.1 add_callback()

```
def codar.savanna.model.Run.add_callback (
    self,
    fn )
```

Function takes single argument which is this run instance, and is called when the process is complete (either normally or killed by timeout). Callbacks must not block.

5.12.2.2 exception()

```
def codar.savanna.model.Run.exception (
    self )
```

True if there was a python exception in the run method. When this is the case, the state of the underlying process is unknown - it may have been started or not.

5.12.2.3 from_data()

```
def codar.savanna.model.Run.from_data (
    cls,
    data )
```

Create Run instance from nested dictionary data structure, e.g. parsed from JSON. The keys 'name', 'exe', 'args' are required, all the other keys are optional and have the same names as the constructor args. Raises KeyError if a required key is missing.

5.12.2.4 get_nodes_used()

```
def codar.savanna.model.Run.get_nodes_used (
    self )
```

Get number of nodes needed to run this app. Requires that the pipeline set_ppn method has been called to set this and tasks_per_node on each run.

5.12.2.5 kill()

```
def codar.savanna.model.Run.kill (
    self )
```

Kill process and cause run thread to complete after the wait returns. If the run is already done, does nothing. If the process is killed, it will mark the state as killed so it can be re-run on workflow restart. Thread safe.

5.12.2.6 killed()

```
def codar.savanna.model.Run.killed (
    self )
```

True if the run is done and the kill method was called. Note that this will `_NOT_` be true if an external kill signal caused the process to exit. Raises ValueError if the run is not complete.

5.12.2.7 succeeded()

```
def codar.savanna.model.Run.succeeded (
    self )
```

True if the run is done, finished normally, and had 0 return value. Raises ValueError if the run is not complete.

5.12.2.8 timed_out()

```
def codar.savanna.model.Run.timed_out (
    self )
```

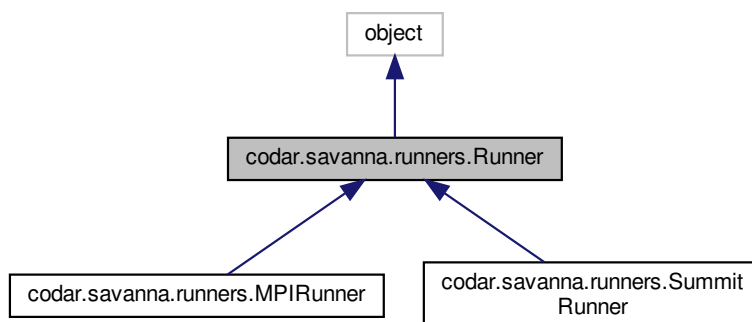
True if the run is done and was killed because it exceeded the specified run timeout. Raises ValueError if the run is not complete.

The documentation for this class was generated from the following file:


- model.py

5.13 codar.savanna.runners.Runner Class Reference

Inheritance diagram for codar.savanna.runners.Runner:



Collaboration diagram for codar.savanna.runners.Runner:



Public Member Functions


- def **wrap** (self, run, sched_args)

The documentation for this class was generated from the following file:


- runners.py

5.14 codar.savanna.exc.SavannaException Class Reference

Inheritance diagram for codar.savanna.exc.SavannaException:



Collaboration diagram for codar.savanna.exc.SavannaException:




The documentation for this class was generated from the following file:


- exc.py

5.15 codar.savanna.machines.SummitNode Class Reference

Inheritance diagram for codar.savanna.machines.SummitNode:



Collaboration diagram for codar.savanna.machines.SummitNode:



Public Member Functions

- def `__init__` (self)
- def `validate_layout` (self)
- def `to_json` (self)

Additional Inherited Members

5.15.1 Member Function Documentation

5.15.1.1 validate_layout()

```
def codar.savanna.machines.SummitNode.validate_layout (
    self )
```


Check that 1) the same rank of the same code is not repeated,
2) a gpu is not mapped to multiple executables.

The documentation for this class was generated from the following file:


- machines.py

5.16 codar.savanna.runners.SummitRunner Class Reference

Inheritance diagram for codar.savanna.runners.SummitRunner:



Collaboration diagram for codar.savanna.runners.SummitRunner:



Public Member Functions

- `def __init__(self)`
- `def wrap(self, run, sched_args)`
- `def wrap_deprecated(self, run, jsrun_opts, find_in_path=True)`

Public Attributes

- `exe`
- `nrs_arg`
- `tasks_per_rs_arg`
- `cpus_per_rs_arg`
- `gpus_per_rs_arg`
- `rs_per_host_arg`
- `launch_distribution_arg`
- `bind_arg`
- `machine`

5.16.1 Member Function Documentation

5.16.1.1 `wrap_deprecated()`

```
def codar.savanna.runners.SummitRunner.wrap_deprecated (
    self,
    run,
    jsrun_opts,
    find_in_path = True )
```


This function is deprecated in favor of the above that uses erf files

The documentation for this class was generated from the following file:


- `runners.py`

5.17 `codar.savanna.status.WorkflowStatus` Class Reference

Inheritance diagram for `codar.savanna.status.WorkflowStatus`:



Collaboration diagram for `codar.savanna.status.WorkflowStatus`:



Public Member Functions

- `def __init__(self, file_path)`
- `def set_state(self, pipeline_state)`

Public Attributes

- `file_path`

The documentation for this class was generated from the following file:

- `status.py`

Index

- `__init__`
 - `codar::savanna::model::NodeConfig`, [20](#)
- `add_callback`
 - `codar::savanna::model::Run`, [32](#)
- `add_node`
 - `codar::savanna::node_layout::NodeLayout`, [22](#)
- `codar.savanna`, [7](#)
- `codar.savanna.consumer`, [7](#)
- `codar.savanna.consumer.PipelineRunner`, [26](#)
- `codar.savanna.exc`, [8](#)
- `codar.savanna.exc.MachineNotFound`, [18](#)
- `codar.savanna.exc.SavannaException`, [35](#)
- `codar.savanna.machines`, [8](#)
- `codar.savanna.machines.Machine`, [16](#)
- `codar.savanna.machines.MachineNode`, [17](#)
- `codar.savanna.machines.SummitNode`, [36](#)
- `codar.savanna.main`, [10](#)
- `codar.savanna.model`, [10](#)
- `codar.savanna.model.NodeConfig`, [20](#)
- `codar.savanna.model.Pipeline`, [24](#)
- `codar.savanna.model.Run`, [30](#)
- `codar.savanna.node_layout.NodeLayout`, [21](#)
- `codar.savanna.producer`, [11](#)
- `codar.savanna.producer.JSONFilePipelineReader`, [15](#)
- `codar.savanna.runners.MPIRunner`, [19](#)
- `codar.savanna.runners.Runner`, [34](#)
- `codar.savanna.runners.SummitRunner`, [37](#)
- `codar.savanna.scheduler`, [11](#)
- `codar.savanna.scheduler.JobList`, [13](#)
- `codar.savanna.status`, [11](#)
- `codar.savanna.status.PipelineState`, [29](#)
- `codar.savanna.status.WorkflowStatus`, [39](#)
- `codar::savanna::consumer::PipelineRunner`
 - `kill_all`, [28](#)
 - `pipeline_finished`, [28](#)
 - `run_finished`, [28](#)
 - `run_pipelines`, [28](#)
 - `stop`, [28](#)
- `codar::savanna::machines`
 - `cori`, [8](#)
 - `summit`, [9](#)
 - `theta`, [9](#)
 - `titan`, [9](#)
- `codar::savanna::machines::Machine`
 - `get_scheduler_options`, [17](#)
- `codar::savanna::machines::SummitNode`
 - `validate_layout`, [36](#)
- `codar::savanna::model::NodeConfig`
 - `__init__`, [20](#)
- `codar::savanna::model::Pipeline`
 - `force_kill_all`, [25](#)
 - `from_data`, [25](#)
 - `set_ppn`, [25](#)
 - `set_total_nodes`, [26](#)
- `codar::savanna::model::Run`
 - `add_callback`, [32](#)
 - `exception`, [32](#)
 - `from_data`, [32](#)
 - `get_nodes_used`, [32](#)
 - `kill`, [33](#)
 - `killed`, [33](#)
 - `succeeded`, [33](#)
 - `timed_out`, [33](#)
- `codar::savanna::node_layout::NodeLayout`
 - `add_node`, [22](#)
 - `default_no_share_layout`, [22](#)
 - `get_node_containing_code`, [22](#)
 - `group_codes_by_node`, [23](#)
 - `serialize_to_dict`, [23](#)
 - `validate`, [23](#)
- `codar::savanna::runners::SummitRunner`
 - `wrap_deprecated`, [38](#)
- `codar::savanna::scheduler::JobList`
 - `pop_job`, [14](#)
- `cori`
 - `codar::savanna::machines`, [8](#)
- `default_no_share_layout`
 - `codar::savanna::node_layout::NodeLayout`, [22](#)
- `exception`
 - `codar::savanna::model::Run`, [32](#)
- `force_kill_all`
 - `codar::savanna::model::Pipeline`, [25](#)
- `from_data`
 - `codar::savanna::model::Pipeline`, [25](#)
 - `codar::savanna::model::Run`, [32](#)
- `get_node_containing_code`
 - `codar::savanna::node_layout::NodeLayout`, [22](#)
- `get_nodes_used`
 - `codar::savanna::model::Run`, [32](#)
- `get_scheduler_options`
 - `codar::savanna::machines::Machine`, [17](#)
- `group_codes_by_node`
 - `codar::savanna::node_layout::NodeLayout`, [23](#)
- `kill`

- `codar::savanna::model::Run`, [33](#)
- `kill_all`
 - `codar::savanna::consumer::PipelineRunner`, [28](#)
- `killed`
 - `codar::savanna::model::Run`, [33](#)
- `pipeline_finished`
 - `codar::savanna::consumer::PipelineRunner`, [28](#)
- `pop_job`
 - `codar::savanna::scheduler::JobList`, [14](#)
- `run_finished`
 - `codar::savanna::consumer::PipelineRunner`, [28](#)
- `run_pipelines`
 - `codar::savanna::consumer::PipelineRunner`, [28](#)
- `serialize_to_dict`
 - `codar::savanna::node_layout::NodeLayout`, [23](#)
- `set_ppn`
 - `codar::savanna::model::Pipeline`, [25](#)
- `set_total_nodes`
 - `codar::savanna::model::Pipeline`, [26](#)
- `stop`
 - `codar::savanna::consumer::PipelineRunner`, [28](#)
- `succeeded`
 - `codar::savanna::model::Run`, [33](#)
- `summit`
 - `codar::savanna::machines`, [9](#)
- `theta`
 - `codar::savanna::machines`, [9](#)
- `timed_out`
 - `codar::savanna::model::Run`, [33](#)
- `titan`
 - `codar::savanna::machines`, [9](#)
- `validate`
 - `codar::savanna::node_layout::NodeLayout`, [23](#)
- `validate_layout`
 - `codar::savanna::machines::SummitNode`, [36](#)
- `wrap_deprecated`
 - `codar::savanna::runners::SummitRunner`, [38](#)